

# CrimeDetection

May 5, 2018

## 1 Problem Description : Crime Detection Regression Analysis

```
In [1]: from keras.models import Sequential
        from keras.layers import Dense
        from keras.layers import LSTM
        from keras.layers import Dropout
        from keras.layers import Flatten, Input
        from keras import backend as K
        from keras.models import Model, load_model
        from sklearn.preprocessing import StandardScaler
        import matplotlib.pyplot as plt
        from sklearn.model_selection import KFold
        from scipy.spatial import distance
        from sklearn.decomposition import PCA
        from numpy import linalg as LA
        from keras.objectives import categorical_crossentropy
        from sklearn.metrics import roc_curve, auc
        import math
        from scipy.stats import pearsonr
        import copy
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error, r2_score
        import itertools
        import csv
        from sklearn import metrics
        import tensorflow as tf
        import tensorflow.contrib.layers as tl
        import numpy as np
        import pandas as pd
        from sklearn import linear_model
        from sklearn.ensemble import RandomForestClassifier
        import seaborn as sb
        %matplotlib inline
```

```
/home/ramchalamkr/.local/lib/python2.7/site-packages/h5py/__init__.py:36: FutureWarning: Conver
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

## 1.1 Steps

### 1.2 1. Checking Data

```
In [18]: X = pd.read_csv('crime_prep.csv',delimiter=',')
         print X.shape
         X[0:20]
```

(1994, 128)

```
Out[18]:
```

	target	v_cont_0	v_cat_0	v_cat_1		v_cat_2	v_cat_3	\
0	0.20	8	NaN	NaN		Lakewoodcity	1	
1	0.67	53	NaN	NaN		Tukwilacity	1	
2	0.43	24	NaN	NaN		Aberdeentown	1	
3	0.12	34	5.0	81440.0	Willingborotownship		1	
4	0.03	42	95.0	6096.0	Bethlehemtownship		1	
5	0.14	6	NaN	NaN	SouthPasadenacity		1	
6	0.03	44	7.0	41500.0	Lincolntown		1	
7	0.55	6	NaN	NaN	Selmacity		1	
8	0.53	21	NaN	NaN	Hendersoncity		1	
9	0.15	29	NaN	NaN	Claytoncity		1	
10	0.24	6	NaN	NaN	DalyCitycity		1	
11	0.08	36	NaN	NaN	RockvilleCentrevillage		1	
12	0.06	25	21.0	44105.0	Needhamtown		1	
13	0.09	55	87.0	30075.0	GrandChutetown		1	
14	0.21	6	NaN	NaN	DanaPointcity		1	
15	0.30	19	187.0	91370.0	FortDodgecity		1	
16	0.49	36	1.0	1000.0	Albanycity		1	
17	0.07	34	27.0	17650.0	Denvilletownship		1	
18	0.15	18	NaN	NaN	Valparaisocity		1	
19	0.03	42	129.0	66376.0	Rostravertownship		1	

	v_cont_5	v_cont_6	v_cont_7	v_cont_8	...	v_cont_117	\
0	0.19	0.33	0.02	0.90	...	0.29	
1	0.00	0.16	0.12	0.74	...	NaN	
2	0.00	0.42	0.49	0.56	...	NaN	
3	0.04	0.77	1.00	0.08	...	NaN	
4	0.01	0.55	0.02	0.95	...	NaN	
5	0.02	0.28	0.06	0.54	...	NaN	
6	0.01	0.39	0.00	0.98	...	NaN	
7	0.01	0.74	0.03	0.46	...	NaN	
8	0.03	0.34	0.20	0.84	...	NaN	
9	0.01	0.40	0.06	0.87	...	NaN	
10	0.13	0.71	0.15	0.07	...	NaN	
11	0.02	0.46	0.08	0.91	...	NaN	
12	0.03	0.47	0.01	0.96	...	NaN	
13	0.01	0.44	0.00	0.98	...	NaN	
14	0.04	0.36	0.01	0.85	...	NaN	

15	0.03	0.34	0.06	0.93	...	NaN
16	0.15	0.31	0.40	0.63	...	0.22
17	0.01	0.53	0.01	0.94	...	NaN
18	0.02	0.47	0.01	0.97	...	NaN
19	0.00	0.41	0.05	0.96	...	NaN

	v_cont_118	v_cont_119	v_cont_120	v_cont_121	v_cont_122	v_cont_123	\
0	0.12	0.26	0.20	0.06	0.04	0.90	
1	0.02	0.12	0.45	NaN	NaN	NaN	
2	0.01	0.21	0.02	NaN	NaN	NaN	
3	0.02	0.39	0.28	NaN	NaN	NaN	
4	0.04	0.09	0.02	NaN	NaN	NaN	
5	0.01	0.58	0.10	NaN	NaN	NaN	
6	0.05	0.08	0.06	NaN	NaN	NaN	
7	0.01	0.33	0.00	NaN	NaN	NaN	
8	0.04	0.17	0.04	NaN	NaN	NaN	
9	0.00	0.47	0.11	NaN	NaN	NaN	
10	0.02	1.00	1.00	NaN	NaN	NaN	
11	0.01	0.63	1.00	NaN	NaN	NaN	
12	0.03	0.18	0.59	NaN	NaN	NaN	
13	0.08	0.04	0.00	NaN	NaN	NaN	
14	0.02	0.40	0.15	NaN	NaN	NaN	
15	0.04	0.15	0.04	NaN	NaN	NaN	
16	0.06	0.39	0.84	0.06	0.06	0.91	
17	0.03	0.09	0.21	NaN	NaN	NaN	
18	0.03	0.20	0.07	NaN	NaN	NaN	
19	0.09	0.03	0.05	NaN	NaN	NaN	

	v_cont_124	v_cont_125	v_cont_126
0	0.5	0.32	0.14
1	NaN	0.00	NaN
2	NaN	0.00	NaN
3	NaN	0.00	NaN
4	NaN	0.00	NaN
5	NaN	0.00	NaN
6	NaN	0.00	NaN
7	NaN	0.00	NaN
8	NaN	0.00	NaN
9	NaN	0.00	NaN
10	NaN	0.00	NaN
11	NaN	0.00	NaN
12	NaN	0.00	NaN
13	NaN	0.00	NaN
14	NaN	0.00	NaN
15	NaN	0.00	NaN
16	0.5	0.88	0.26
17	NaN	0.00	NaN
18	NaN	0.00	NaN

19                NaN                0.00                NaN

[20 rows x 128 columns]

### 1.3 2. Tidying the data

Missing data can either be filled with the means of the features or maybe 0 to ignore them or even other data imputation techniques to predict the missing values

In this case the mean of the features is taken for the missing values.

```
In [19]: print len(X['v_cat_2'].unique())
X = X.fillna(X.mean())
del X['v_cat_2']
del X['v_cat_3']
Y = X['target']
del X['target']
X.head()
```

1828

```
Out[19]:
```

	v_cont_0	v_cat_0	v_cat_1	v_cont_5	v_cont_6	v_cont_7	v_cont_8	\
0	8	58.826829	46188.336597	0.19	0.33	0.02	0.90	
1	53	58.826829	46188.336597	0.00	0.16	0.12	0.74	
2	24	58.826829	46188.336597	0.00	0.42	0.49	0.56	
3	34	5.000000	81440.000000	0.04	0.77	1.00	0.08	
4	42	95.000000	6096.000000	0.01	0.55	0.02	0.95	

	v_cont_9	v_cont_10	v_cont_11	...	v_cont_117	v_cont_118	\
0	0.12	0.17	0.34	...	0.290000	0.12	
1	0.45	0.07	0.26	...	0.305987	0.02	
2	0.17	0.04	0.39	...	0.305987	0.01	
3	0.12	0.10	0.51	...	0.305987	0.02	
4	0.09	0.05	0.38	...	0.305987	0.04	

	v_cont_119	v_cont_120	v_cont_121	v_cont_122	v_cont_123	v_cont_124	\
0	0.26	0.20	0.060000	0.040000	0.900000	0.500000	
1	0.12	0.45	0.163103	0.076708	0.698589	0.440439	
2	0.21	0.02	0.163103	0.076708	0.698589	0.440439	
3	0.39	0.28	0.163103	0.076708	0.698589	0.440439	
4	0.09	0.02	0.163103	0.076708	0.698589	0.440439	

	v_cont_125	v_cont_126
0	0.32	0.140000
1	0.00	0.195078
2	0.00	0.195078
3	0.00	0.195078

```
4          0.00      0.195078
```

```
[5 rows x 125 columns]
```

```
In [20]: #Standardise certain columns
```

```
X['v_cat_0'] = StandardScaler().fit_transform(X['v_cat_0'].values.reshape(-1, 1))
X['v_cat_1'] = StandardScaler().fit_transform(X['v_cat_1'].values.reshape(-1, 1))
X['v_cont_0'] = StandardScaler().fit_transform(X['v_cont_0'].values.reshape(-1, 1))
X.head()
```

```
/usr/local/lib/python2.7/dist-packages/sklearn/utils/validation.py:444: DataConversionWarning:
  warnings.warn(msg, DataConversionWarning)
```

```
Out [20]:
```

	v_cont_0	v_cat_0	v_cat_1	v_cont_5	v_cont_6	v_cont_7	v_cont_8	\
0	-1.261697	0.000000	0.000000	0.19	0.33	0.02	0.90	
1	1.483304	0.000000	0.000000	0.00	0.16	0.12	0.74	
2	-0.285697	0.000000	0.000000	0.00	0.42	0.49	0.56	
3	0.324303	-0.664357	2.178117	0.04	0.77	1.00	0.08	
4	0.812304	0.446467	-2.477211	0.01	0.55	0.02	0.95	

	v_cont_9	v_cont_10	v_cont_11	...	v_cont_117	v_cont_118	\
0	0.12	0.17	0.34	...	0.290000	0.12	
1	0.45	0.07	0.26	...	0.305987	0.02	
2	0.17	0.04	0.39	...	0.305987	0.01	
3	0.12	0.10	0.51	...	0.305987	0.02	
4	0.09	0.05	0.38	...	0.305987	0.04	

	v_cont_119	v_cont_120	v_cont_121	v_cont_122	v_cont_123	v_cont_124	\
0	0.26	0.20	0.060000	0.040000	0.900000	0.500000	
1	0.12	0.45	0.163103	0.076708	0.698589	0.440439	
2	0.21	0.02	0.163103	0.076708	0.698589	0.440439	
3	0.39	0.28	0.163103	0.076708	0.698589	0.440439	
4	0.09	0.02	0.163103	0.076708	0.698589	0.440439	

	v_cont_125	v_cont_126
0	0.32	0.140000
1	0.00	0.195078
2	0.00	0.195078
3	0.00	0.195078
4	0.00	0.195078

```
[5 rows x 125 columns]
```

## 1.4 3. Exploratory Analysis

### 1.4.1 PCA

```
In [21]: print X.shape
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3,random_state=42)
```

```

Covariance = np.dot(X_train.T,X_train)
print Covariance.shape

Lambda, e = LA.eigh(Covariance)
Lambda = Lambda.reshape(Lambda.shape[0],1)
Lambda = sorted(Lambda,reverse=True)
TotalLambda = np.sum(Lambda)
LambdaProp = []
for i in range(X_train.shape[1]):
    temp = np.sum(Lambda[0:i])*1.0/TotalLambda
    LambdaProp.append(temp)


Dim = np.linspace(0, X_train.shape[1], X_train.shape[1])

plt.plot(Dim,LambdaProp)
plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
          ncol=2, mode="expand", borderaxespad=0.)
plt.xlabel('Dimensions')
plt.ylabel('Lambda/TotalLambda')
plt.show()

```

(1994, 125)

(125, 125)



```

In [23]: X1 = X
         Y1 = Y

         for i in (range(1,100)):
             print i
             NewFeatureSet = np.dot(X1,e[:,125-i:125])
             print NewFeatureSet.shape
             X_train,X_test,Y_train,Y_test = train_test_split(NewFeatureSet,Y1,test_size = 0.2)
             #Normal Linear Regression
             regr = linear_model.LinearRegression(normalize = True)
             regr.fit(X_train,Y_train)
             Y_pred = regr.predict(X_test)
             Weights = regr.coef_
             print "RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred))
             print('Variance/R2 score: %.2f' % r2_score(Y_test, Y_pred))

1
(1994, 1)
RMSE 0.22
Variance/R2 score: -0.00
2
(1994, 2)
RMSE 0.21
Variance/R2 score: 0.04
3
(1994, 3)
RMSE 0.17
Variance/R2 score: 0.39
4
(1994, 4)
RMSE 0.16
Variance/R2 score: 0.48
5
(1994, 5)
RMSE 0.15
Variance/R2 score: 0.52
6
(1994, 6)
RMSE 0.15
Variance/R2 score: 0.52
7
(1994, 7)
RMSE 0.15
Variance/R2 score: 0.54
8
(1994, 8)
RMSE 0.15
Variance/R2 score: 0.55

```

9  
(1994, 9)  
RMSE 0.14  
Variance/R2 score: 0.59  
10  
(1994, 10)  
RMSE 0.14  
Variance/R2 score: 0.61  
11  
(1994, 11)  
RMSE 0.14  
Variance/R2 score: 0.61  
12  
(1994, 12)  
RMSE 0.13  
Variance/R2 score: 0.62  
13  
(1994, 13)  
RMSE 0.13  
Variance/R2 score: 0.63  
14  
(1994, 14)  
RMSE 0.13  
Variance/R2 score: 0.63  
15  
(1994, 15)  
RMSE 0.13  
Variance/R2 score: 0.63  
16  
(1994, 16)  
RMSE 0.13  
Variance/R2 score: 0.63  
17  
(1994, 17)  
RMSE 0.13  
Variance/R2 score: 0.63  
18  
(1994, 18)  
RMSE 0.13  
Variance/R2 score: 0.64  
19  
(1994, 19)  
RMSE 0.13  
Variance/R2 score: 0.64  
20  
(1994, 20)  
RMSE 0.13  
Variance/R2 score: 0.64



21  
(1994, 21)  
RMSE 0.13  
Variance/R2 score: 0.64  
22  
(1994, 22)  
RMSE 0.13  
Variance/R2 score: 0.64  
23  
(1994, 23)  
RMSE 0.13  
Variance/R2 score: 0.64  
24  
(1994, 24)  
RMSE 0.13  
Variance/R2 score: 0.64  
25  
(1994, 25)  
RMSE 0.13  
Variance/R2 score: 0.64  
26  
(1994, 26)  
RMSE 0.13  
Variance/R2 score: 0.64  
27  
(1994, 27)  
RMSE 0.13  
Variance/R2 score: 0.64  
28  
(1994, 28)  
RMSE 0.13  
Variance/R2 score: 0.65  
29  
(1994, 29)  
RMSE 0.13  
Variance/R2 score: 0.64  
30  
(1994, 30)  
RMSE 0.13  
Variance/R2 score: 0.64  
31  
(1994, 31)  
RMSE 0.13  
Variance/R2 score: 0.64  
32  
(1994, 32)  
RMSE 0.13  
Variance/R2 score: 0.64

33  
(1994, 33)  
RMSE 0.13  
Variance/R2 score: 0.64  
34  
(1994, 34)  
RMSE 0.13  
Variance/R2 score: 0.64  
35  
(1994, 35)  
RMSE 0.13  
Variance/R2 score: 0.64  
36  
(1994, 36)  
RMSE 0.13  
Variance/R2 score: 0.64  
37  
(1994, 37)  
RMSE 0.13  
Variance/R2 score: 0.64  
38  
(1994, 38)  
RMSE 0.13  
Variance/R2 score: 0.63  
39  
(1994, 39)  
RMSE 0.13  
Variance/R2 score: 0.63  
40  
(1994, 40)  
RMSE 0.13  
Variance/R2 score: 0.64  
41  
(1994, 41)  
RMSE 0.13  
Variance/R2 score: 0.63  
42  
(1994, 42)  
RMSE 0.13  
Variance/R2 score: 0.64  
43  
(1994, 43)  
RMSE 0.13  
Variance/R2 score: 0.64  
44  
(1994, 44)  
RMSE 0.13  
Variance/R2 score: 0.64

45  
(1994, 45)  
RMSE 0.13  
Variance/R2 score: 0.64  
46  
(1994, 46)  
RMSE 0.13  
Variance/R2 score: 0.64  
47  
(1994, 47)  
RMSE 0.13  
Variance/R2 score: 0.64  
48  
(1994, 48)  
RMSE 0.13  
Variance/R2 score: 0.63  
49  
(1994, 49)  
RMSE 0.13  
Variance/R2 score: 0.63  
50  
(1994, 50)  
RMSE 0.13  
Variance/R2 score: 0.63  
51  
(1994, 51)  
RMSE 0.13  
Variance/R2 score: 0.63  
52  
(1994, 52)  
RMSE 0.13  
Variance/R2 score: 0.63  
53  
(1994, 53)  
RMSE 0.13  
Variance/R2 score: 0.63  
54  
(1994, 54)  
RMSE 0.13  
Variance/R2 score: 0.62  
55  
(1994, 55)  
RMSE 0.13  
Variance/R2 score: 0.62  
56  
(1994, 56)  
RMSE 0.14  
Variance/R2 score: 0.62

57  
(1994, 57)  
RMSE 0.14  
Variance/R2 score: 0.62  
58  
(1994, 58)  
RMSE 0.14  
Variance/R2 score: 0.62  
59  
(1994, 59)  
RMSE 0.14  
Variance/R2 score: 0.62  
60  
(1994, 60)  
RMSE 0.13  
Variance/R2 score: 0.62  
61  
(1994, 61)  
RMSE 0.13  
Variance/R2 score: 0.62  
62  
(1994, 62)  
RMSE 0.13  
Variance/R2 score: 0.62  
63  
(1994, 63)  
RMSE 0.13  
Variance/R2 score: 0.62  
64  
(1994, 64)  
RMSE 0.14  
Variance/R2 score: 0.62  
65  
(1994, 65)  
RMSE 0.13  
Variance/R2 score: 0.62  
66  
(1994, 66)  
RMSE 0.14  
Variance/R2 score: 0.62  
67  
(1994, 67)  
RMSE 0.14  
Variance/R2 score: 0.62  
68  
(1994, 68)  
RMSE 0.14  
Variance/R2 score: 0.62

69  
(1994, 69)  
RMSE 0.14  
Variance/R2 score: 0.62  
70  
(1994, 70)  
RMSE 0.14  
Variance/R2 score: 0.62  
71  
(1994, 71)  
RMSE 0.14  
Variance/R2 score: 0.62  
72  
(1994, 72)  
RMSE 0.14  
Variance/R2 score: 0.61  
73  
(1994, 73)  
RMSE 0.14  
Variance/R2 score: 0.61  
74  
(1994, 74)  
RMSE 0.14  
Variance/R2 score: 0.61  
75  
(1994, 75)  
RMSE 0.14  
Variance/R2 score: 0.61  
76  
(1994, 76)  
RMSE 0.14  
Variance/R2 score: 0.61  
77  
(1994, 77)  
RMSE 0.14  
Variance/R2 score: 0.61  
78  
(1994, 78)  
RMSE 0.14  
Variance/R2 score: 0.61  
79  
(1994, 79)  
RMSE 0.14  
Variance/R2 score: 0.61  
80  
(1994, 80)  
RMSE 0.14  
Variance/R2 score: 0.61

81  
(1994, 81)  
RMSE 0.14  
Variance/R2 score: 0.61  
82  
(1994, 82)  
RMSE 0.14  
Variance/R2 score: 0.61  
83  
(1994, 83)  
RMSE 0.14  
Variance/R2 score: 0.61  
84  
(1994, 84)  
RMSE 0.14  
Variance/R2 score: 0.61  
85  
(1994, 85)  
RMSE 0.14  
Variance/R2 score: 0.61  
86  
(1994, 86)  
RMSE 0.14  
Variance/R2 score: 0.61  
87  
(1994, 87)  
RMSE 0.14  
Variance/R2 score: 0.61  
88  
(1994, 88)  
RMSE 0.14  
Variance/R2 score: 0.61  
89  
(1994, 89)  
RMSE 0.14  
Variance/R2 score: 0.60  
90  
(1994, 90)  
RMSE 0.14  
Variance/R2 score: 0.61  
91  
(1994, 91)  
RMSE 0.14  
Variance/R2 score: 0.61  
92  
(1994, 92)  
RMSE 0.14  
Variance/R2 score: 0.61

```

93
(1994, 93)
RMSE 0.14
Variance/R2 score: 0.61
94
(1994, 94)
RMSE 0.14
Variance/R2 score: 0.61
95
(1994, 95)
RMSE 0.14
Variance/R2 score: 0.61
96
(1994, 96)
RMSE 0.14
Variance/R2 score: 0.62
97
(1994, 97)
RMSE 0.14
Variance/R2 score: 0.62
98
(1994, 98)
RMSE 0.14
Variance/R2 score: 0.62
99
(1994, 99)
RMSE 0.14
Variance/R2 score: 0.62

```

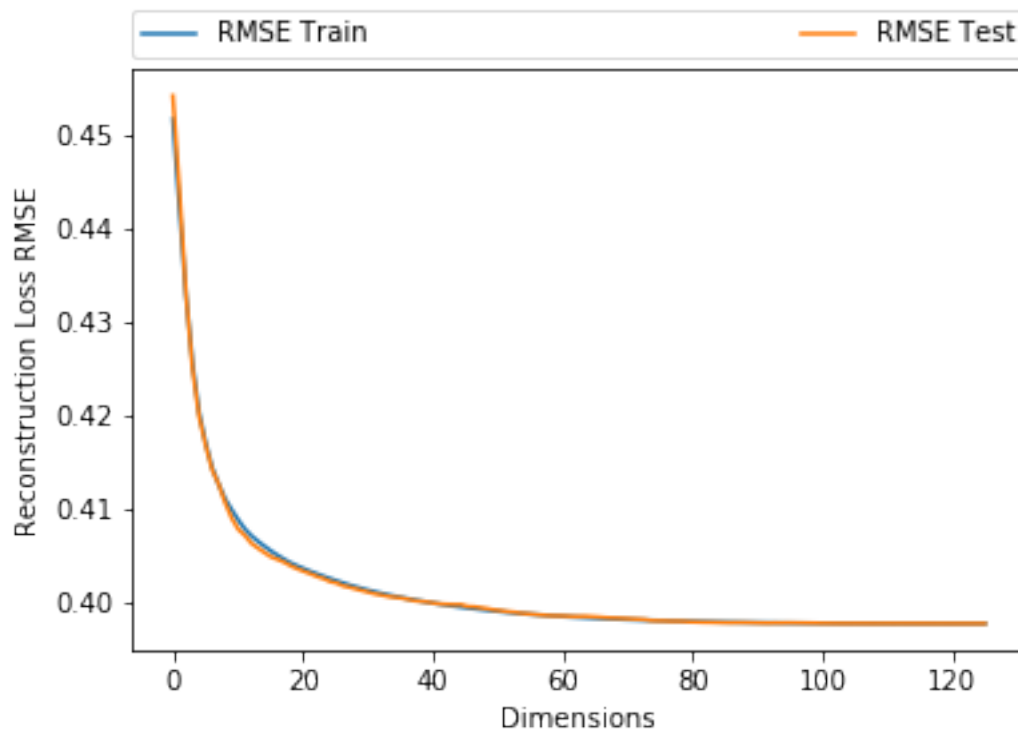
```

In [24]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3,random_state=42)
        RMSETrain = []
        RMSETest = []
        VarainceRatio = []
        Dim = np.linspace(0, X_train.shape[1], X_train.shape[1])
        for i in range(1,X_train.shape[1]+1):
            pca = PCA(n_components = i)
            pca.fit(X_train)
            VarainceRatio.append(pca.explained_variance_ratio_)
            X_trainhat = np.dot(pca.transform(X_train)[:,:i], pca.components_[i,:])
            RMSETrain.append(math.sqrt(mean_squared_error(X_train,X_trainhat)))
            X_testthat = np.dot(pca.transform(X_test)[:,:i], pca.components_[i,:])
            RMSETest.append(math.sqrt(mean_squared_error(X_test,X_testthat)))

        plt.plot(Dim,RMSETrain,label="RMSE Train")
        plt.plot(Dim,RMSETest,label="RMSE Test")
        plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
                    ncol=2, mode="expand", borderaxespad=0.)

```

```
plt.xlabel('Dimensions')
plt.ylabel('Reconstruction Loss RMSE')
plt.show()
```



```
In [25]: X1 = X
        Y1 = Y
        for i in range(1,15):
            pca = PCA(n_components = i)
            X_1 = pca.fit_transform(X1)
            #VarainceRatio.append(pca.explained_variance_ratio_)
            print X_1
            print "no of componenets",i
            print "predict for transformed data"

            X_train,X_test,Y_train,Y_test = train_test_split(X_1,Y1,test_size = 0.2,random_state=42)
            #Normal Linear Regression
            regr = linear_model.LinearRegression(normalize = True)
            regr.fit(X_train,Y_train)
            Y_pred = regr.predict(X_test)
            Weights = regr.coef_
            print "RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred))
            print('Variance/R2 score: %.2f' % r2_score(Y_test, Y_pred))
```



```

[[-0.83893174]
 [ 0.91684945]
 [ 0.21196309]
 ...
 [-0.6326288 ]
 [-0.63799423]
 [-1.45620074]]
no of componenets 1
predict for transformed data
RMSE 0.22
Variance/R2 score: -0.02
[[-0.83893174  0.19106065]
 [ 0.91684946 -0.35978791]
 [ 0.21196309  0.41829515]
 ...
 [-0.6326288  0.70376198]
 [-0.63799423 -0.32252535]
 [-1.45620074  1.73787881]]
no of componenets 2
predict for transformed data
RMSE 0.16
Variance/R2 score: 0.49
[[-0.83893174  0.19106065  0.39788124]
 [ 0.91684946 -0.35978791 -0.32442181]
 [ 0.21196309  0.41829515 -0.13784697]
 ...
 [-0.6326288  0.70376198  1.6834653 ]
 [-0.63799423 -0.32252534  1.46103493]
 [-1.45620074  1.73787882  0.58259365]]
no of componenets 3
predict for transformed data
RMSE 0.16
Variance/R2 score: 0.48
[[-0.83893174  0.19106065  0.39788124  0.28589089]
 [ 0.91684946 -0.35978791 -0.32442181 -0.0468139 ]
 [ 0.21196309  0.41829515 -0.13784697 -0.30614276]
 ...
 [-0.6326288  0.70376198  1.6834653 -1.63788755]
 [-0.63799423 -0.32252534  1.46103493 -0.80879764]
 [-1.45620074  1.73787882  0.58259365  0.30865221]]
no of componenets 4
predict for transformed data
RMSE 0.15
Variance/R2 score: 0.50
[[-0.83893174  0.19106065  0.39788124  0.28589089 -0.9762273 ]
 [ 0.91684946 -0.35978791 -0.32442181 -0.0468139  0.86189957]
 [ 0.21196309  0.41829515 -0.13784697 -0.30614276 -0.70036877]
 ...

```

```

[-0.6326288  0.70376198  1.6834653  -1.63788755 -0.3569823 ]
[-0.63799423 -0.32252534  1.46103493 -0.80879764  0.67782456]
[-1.45620074  1.73787882  0.58259365  0.30865221  1.00820566]]
no of componenets 5
predict for transformed data
RMSE 0.15
Variance/R2 score: 0.51
[[-0.83893174  0.19106065  0.39788124  0.28589089 -0.97622731  0.38611268]
 [ 0.91684946 -0.35978791 -0.32442181 -0.0468139  0.86189956  1.07084359]
 [ 0.21196309  0.41829515 -0.13784697 -0.30614276 -0.70036876 -0.16831285]
 ...
 [-0.6326288  0.70376198  1.6834653  -1.63788755 -0.35698233  0.25897959]
 [-0.63799423 -0.32252534  1.46103493 -0.80879764  0.67782455  0.4963086 ]
 [-1.45620074  1.73787882  0.58259365  0.30865221  1.00820565 -0.71288661]]
no of componenets 6
predict for transformed data
RMSE 0.15
Variance/R2 score: 0.51
[[-8.38931741e-01  1.91060648e-01  3.97881242e-01 ... -9.76227318e-01
  3.86112475e-01  1.63136503e-01]
 [ 9.16849455e-01 -3.59787912e-01 -3.24421808e-01 ...  8.61899550e-01
  1.07084340e+00  2.51164529e-01]
 [ 2.11963089e-01  4.18295154e-01 -1.37846971e-01 ... -7.00368786e-01
 -1.68313217e-01  2.32590950e-04]
 ...
 [-6.32628800e-01  7.03761983e-01  1.68346530e+00 ... -3.56982307e-01
  2.58979854e-01 -8.13425272e-01]
 [-6.37994231e-01 -3.22525345e-01  1.46103493e+00 ...  6.77824565e-01
  4.96308786e-01  6.61473722e-02]
 [-1.45620074e+00  1.73787882e+00  5.82593653e-01 ...  1.00820566e+00
 -7.12886689e-01  1.93132382e-01]]
no of componenets 7
predict for transformed data
RMSE 0.15
Variance/R2 score: 0.55
[[-8.38931741e-01  1.91060648e-01  3.97881242e-01 ...  3.86112724e-01
  1.63136417e-01 -5.09189088e-01]
 [ 9.16849455e-01 -3.59787912e-01 -3.24421808e-01 ...  1.07084352e+00
  2.51164623e-01 -1.49584591e-01]
 [ 2.11963089e-01  4.18295154e-01 -1.37846971e-01 ... -1.68312990e-01
  2.32076835e-04 -4.38400635e-02]
 ...
 [-6.32628800e-01  7.03761983e-01  1.68346530e+00 ...  2.58979684e-01
 -8.13425816e-01 -1.68552412e-01]
 [-6.37994231e-01 -3.22525345e-01  1.46103493e+00 ...  4.96308677e-01
  6.61473485e-02  2.97791972e-01]
 [-1.45620074e+00  1.73787882e+00  5.82593653e-01 ... -7.12886814e-01
  1.93131924e-01 -5.57367450e-01]]

```

```

no of componenets 8
predict for transformed data
RMSE 0.14
Variance/R2 score: 0.59
[[-8.38931741e-01  1.91060648e-01  3.97881242e-01 ...  1.63136521e-01
  -5.09186136e-01 -4.34768594e-01]
 [ 9.16849455e-01 -3.59787912e-01 -3.24421808e-01 ...  2.51164624e-01
  -1.49586134e-01 -8.24804485e-01]
 [ 2.11963089e-01  4.18295154e-01 -1.37846971e-01 ...  2.32085531e-04
  -4.38406691e-02  1.48838088e-01]
 ...
 [-6.32628800e-01  7.03761983e-01  1.68346530e+00 ... -8.13425763e-01
  -1.68551238e-01 -6.53347688e-02]
 [-6.37994231e-01 -3.22525345e-01  1.46103493e+00 ...  6.61473512e-02
  2.97790406e-01 -1.26347619e-01]
 [-1.45620074e+00  1.73787882e+00  5.82593653e-01 ...  1.93132012e-01
  -5.57363522e-01 -3.25756539e-01]]

```

```

no of componenets 9
predict for transformed data
RMSE 0.14
Variance/R2 score: 0.61
[[-0.83893174  0.19106065  0.39788124 ... -0.50918594 -0.43477104
  0.04578273]
 [ 0.91684946 -0.35978791 -0.32442181 ... -0.14958673 -0.824802
  0.24931816]
 [ 0.21196309  0.41829515 -0.13784697 ... -0.04383982  0.14883761
  0.44705059]
 ...
 [-0.6326288  0.70376198  1.6834653 ... -0.16855119 -0.06534024
  -0.13388745]
 [-0.63799423 -0.32252534  1.46103493 ...  0.29779086 -0.12634794
  -0.58565302]
 [-1.45620074  1.73787882  0.58259365 ... -0.55736359 -0.32576343
  0.00474748]]

```

```

no of componenets 10
predict for transformed data
RMSE 0.13
Variance/R2 score: 0.63
[[-0.83893174  0.19106065  0.39788124 ... -0.43476782  0.0457842
  0.02812753]
 [ 0.91684946 -0.35978791 -0.32442181 ... -0.82480229  0.24931289
  -0.72522575]
 [ 0.21196309  0.41829515 -0.13784697 ...  0.14883595  0.44705413
  -0.50833336]
 ...
 [-0.6326288  0.70376198  1.6834653 ... -0.0653433 -0.13387715
  0.1310201 ]
 [-0.63799423 -0.32252534  1.46103493 ... -0.12635255 -0.58564017

```

```

-0.0699597 ]
[-1.45620074  1.73787882  0.58259365 ... -0.3257647  0.00472953
 0.03595945]]
no of componenets 11
predict for transformed data
RMSE 0.13
Variance/R2 score: 0.63
[[-0.83893174  0.19106065  0.39788124 ...  0.04575498  0.02812399
  0.18395406]
 [ 0.91684946 -0.35978791 -0.32442181 ...  0.24932038 -0.72522364
 -0.02691651]
 [ 0.21196309  0.41829515 -0.13784697 ...  0.447057  -0.50832742
 -0.01767973]
 ...
 [-0.6326288  0.70376198  1.6834653 ... -0.13384475  0.13102212
  0.10336862]
 [-0.63799423 -0.32252534  1.46103493 ... -0.58561298 -0.06995331
 -0.0472847 ]
 [-1.45620074  1.73787882  0.58259365 ...  0.00476318  0.03596686
  0.10974159]]
no of componenets 12
predict for transformed data
RMSE 0.13
Variance/R2 score: 0.63
[[-0.83893174  0.19106071  0.39788123 ...  0.02780265  0.18346462
  0.08048317]
 [ 0.91684945 -0.35978793 -0.32442186 ... -0.725408  -0.0270232
 -0.15179718]
 [ 0.2119631  0.41829514 -0.13784691 ... -0.50811985 -0.01690639
  0.38017355]
 ...
 [-0.6326288  0.70376196  1.68346533 ...  0.13085839  0.10375661
  0.44058922]
 [-0.63799422 -0.32252537  1.46103492 ... -0.0696856  -0.04660295
  0.51356321]
 [-1.45620074  1.7378788  0.58259357 ...  0.03529639  0.11024077
 -0.0257058 ]]
no of componenets 13
predict for transformed data
RMSE 0.13
Variance/R2 score: 0.63
[[-0.83893171  0.19106063  0.39788118 ...  0.18225289  0.08196732
 -0.22805535]
 [ 0.91684944 -0.3597879  -0.32442179 ... -0.02744453 -0.14949181
 -0.30509614]
 [ 0.21196309  0.41829516 -0.13784697 ... -0.01795268  0.38026324
 -0.27291182]
 ...

```

```

[-0.6326288  0.70376199  1.6834653  ...  0.10205886  0.44238049
 0.05283028]
[-0.63799422 -0.32252531  1.46103489 ... -0.04973349  0.51472393
 0.08188874]
[-1.45620075  1.73787882  0.58259369 ...  0.11085379 -0.02437587
 0.17981367]]
no of componenets 14
predict for transformed data
RMSE 0.13
Variance/R2 score: 0.63

```


### 1.4.2 Feature engineering using Pearson's Coefficient

```

In [26]: X = (X - X.mean())
         Matrix = X.as_matrix()
         #print TrainX['ALSQM_Count']
         r=[]
         p=[]
         for i in range(X.shape[1]):
             t1,t2 = pearsonr(Matrix[:,i],Y)
             #print t1,i,
             r.append((t1,i,X.columns[i],t2))
         r.sort()

In [27]: frequencies = []
         labels = []
         count=0
         for ind,x in enumerate(r):
             if(x[0]<-0.4 or x[0]>0.4):
                 frequencies.append(x[0])
                 labels.append(x[2])
             else:
                 if count<10:
                     labels.append('-')
                     frequencies.append(0.01)
                     count+=1
         print len(frequencies)
         freq_series = pd.Series.from_array(frequencies)
         x_labels = range(len(freq_series))
         plt.figure(figsize=(15, 10))
         ax = freq_series.plot(kind='bar')
         ax.set_title('Correlation between Dependent and Independent variables')
         ax.set_xlabel('Features Columns')
         ax.set_ylabel('Pearson Correlation Score')
         rects = ax.patches
         ax.set_xticklabels(labels)
         plt.show()

```




```
In [28]: IndList = []
         for ind,x in enumerate(r):
             if(x[0]<-0.4 or x[0]>0.4):
                 IndList.append(x[1])
         CorrList = []
         CorrMatrix = np.zeros((len(IndList),len(IndList)))
         Cols = []
         for i in range(len(IndList)):
             Cols.append(X.columns[IndList[i]])
             for j in range(len(IndList)):
                 t1,t2 = pearsonr(Matrix[:,IndList[i]],Matrix[:,IndList[j]])
                 CorrList.append((t1,IndList[i],X.columns[IndList[i]],IndList[j],X.columns[IndList[j]]))
                 CorrMatrix[i][j] = t1
         print CorrMatrix.shape
```

(27, 27)

```
In [29]: fig, ax = plt.subplots(figsize=(20,20))
         sb.heatmap(CorrMatrix,
```

```
xticklabels=Cols,
yticklabels=Cols,linewidths=1,annot=True, ax=ax,square =True)
```

```
plt.show()
```



## 1.5 4. Model Development and Performance

```
In [30]: print IndList
t=[]
for i in IndList:
    t.append(X.columns[i])
print t
FinalTrainSet = np.zeros((len(IndList),Matrix.shape[0]))
```

```

for ind,i in enumerate(IndList):
    FinalTrainSet[ind] = (Matrix[:,i])
FinalTrainSet = FinalTrainSet.T
print FinalTrainSet.shape
print Y.shape

```

```

[47, 46, 6, 48, 49, 18, 70, 76, 22, 15, 32, 74, 30, 71, 52, 72, 77, 33, 80, 35, 31, 41, 44, 43,
['v_cont_49', 'v_cont_48', 'v_cont_8', 'v_cont_50', 'v_cont_51', 'v_cont_20', 'v_cont_72', 'v_
(1994, 27)
(1994,)

```

### 1.5.1 Linear Regression

```

In [31]: X_train,X_test,Y_train,Y_test = train_test_split(FinalTrainSet,Y,test_size = 0.3,rand
#Normal Linear Regression
clf = linear_model.LinearRegression(normalize = True)
clf.fit(X_train,Y_train)
Y_pred = clf.predict(X_test)
Weights = clf.coef_
print Weights
print "RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred))
print('Variance/R2 score: %.2f' % r2_score(Y_test, Y_pred))

# Linear Regression With Ridge regularisation
clf = linear_model.Ridge(alpha = 0.1,normalize = True)
clf.fit(X_train,Y_train)
Y_pred = clf.predict(X_test)
print "RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred))
print('Variance/R2 score: %.2f' % r2_score(Y_test, Y_pred))

# Linear Regression With Lasso regularisation
clf = linear_model.Lasso(alpha = 0.1,normalize = True)
clf.fit(X_train,Y_train)
Y_pred = clf.predict(X_test)
print "RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred))
print('Variance/R2 score: %.2f' % r2_score(Y_test, Y_pred))

[-0.06986744 -0.20319033  0.00282052 -0.0764294  -0.00365693 -0.11640938
 -0.3154949   0.35358964  0.21298279 -0.04158515 -0.16352962  0.26255075
  0.12639153  0.2220687  -0.23364425  0.10495532  0.04729825  0.11573447
  0.02870576 -0.05019055 -0.06960878  0.4324694  -0.61741931  0.26619755
  0.05778619  0.19690957  0.16119416]
RMSE 0.13
Variance/R2 score: 0.64
RMSE 0.13
Variance/R2 score: 0.64

```



RMSE 0.22  
Variance/R2 score: -0.01

## 1.5.2 Neural Nets

```
In [32]: X_train,X_test,Y_train,Y_test = train_test_split(FinalTrainSet,Y,test_size = 0.3,random_state=42)
InputWidth = X_train.shape[1]
K.clear_session()
model = Sequential()
model.add(Dense(128, input_shape = (InputWidth,), activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(1, activation='relu'))
#model.add(Dense(InputWidth, activation='relu'))
print model.summary()
# Compile model
model.compile(loss='mean_squared_error', optimizer='sgd', metrics=['accuracy'])
op = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), nb_epoch=150, batch_size=32)
y_pred = model.predict(X_test)
print "RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred))
print('Variance score: %.2f' % r2_score(Y_test, Y_pred))
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 128)	3584
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 64)	8256
dropout_2 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 32)	2080
dropout_3 (Dropout)	(None, 32)	0
dense_4 (Dense)	(None, 1)	33

Total params: 13,953  
Trainable params: 13,953  
Non-trainable params: 0

```
/usr/local/lib/python2.7/dist-packages/keras/models.py:848: UserWarning: The `nb_epoch` argument
  warnings.warn('The `nb_epoch` argument in `fit` '
```

Train on 1395 samples, validate on 599 samples

Epoch 1/150

14s - loss: 0.0808 - acc: 0.0036 - val\_loss: 0.0559 - val\_acc: 0.0083

Epoch 2/150

0s - loss: 0.0626 - acc: 0.0050 - val\_loss: 0.0459 - val\_acc: 0.0083

Epoch 3/150

0s - loss: 0.0555 - acc: 0.0065 - val\_loss: 0.0394 - val\_acc: 0.0083

Epoch 4/150

0s - loss: 0.0497 - acc: 0.0079 - val\_loss: 0.0347 - val\_acc: 0.0100

Epoch 5/150

0s - loss: 0.0455 - acc: 0.0115 - val\_loss: 0.0313 - val\_acc: 0.0100

Epoch 6/150

0s - loss: 0.0406 - acc: 0.0122 - val\_loss: 0.0288 - val\_acc: 0.0117

Epoch 7/150

0s - loss: 0.0412 - acc: 0.0136 - val\_loss: 0.0268 - val\_acc: 0.0117

Epoch 8/150

0s - loss: 0.0388 - acc: 0.0143 - val\_loss: 0.0255 - val\_acc: 0.0117

Epoch 9/150

0s - loss: 0.0375 - acc: 0.0158 - val\_loss: 0.0242 - val\_acc: 0.0117

Epoch 10/150

0s - loss: 0.0362 - acc: 0.0151 - val\_loss: 0.0232 - val\_acc: 0.0150

Epoch 11/150

0s - loss: 0.0370 - acc: 0.0151 - val\_loss: 0.0226 - val\_acc: 0.0150

Epoch 12/150

0s - loss: 0.0334 - acc: 0.0179 - val\_loss: 0.0219 - val\_acc: 0.0150

Epoch 13/150

0s - loss: 0.0331 - acc: 0.0194 - val\_loss: 0.0215 - val\_acc: 0.0150

Epoch 14/150

0s - loss: 0.0318 - acc: 0.0172 - val\_loss: 0.0210 - val\_acc: 0.0150

Epoch 15/150

0s - loss: 0.0324 - acc: 0.0186 - val\_loss: 0.0207 - val\_acc: 0.0167

Epoch 16/150

0s - loss: 0.0316 - acc: 0.0186 - val\_loss: 0.0205 - val\_acc: 0.0167

Epoch 17/150

0s - loss: 0.0322 - acc: 0.0215 - val\_loss: 0.0203 - val\_acc: 0.0167

Epoch 18/150

0s - loss: 0.0305 - acc: 0.0208 - val\_loss: 0.0201 - val\_acc: 0.0167

Epoch 19/150

0s - loss: 0.0306 - acc: 0.0208 - val\_loss: 0.0200 - val\_acc: 0.0167

Epoch 20/150

0s - loss: 0.0297 - acc: 0.0179 - val\_loss: 0.0199 - val\_acc: 0.0167

Epoch 21/150

0s - loss: 0.0318 - acc: 0.0194 - val\_loss: 0.0198 - val\_acc: 0.0167  
Epoch 22/150  
0s - loss: 0.0318 - acc: 0.0186 - val\_loss: 0.0198 - val\_acc: 0.0167  
Epoch 23/150  
0s - loss: 0.0308 - acc: 0.0201 - val\_loss: 0.0196 - val\_acc: 0.0184  
Epoch 24/150  
0s - loss: 0.0311 - acc: 0.0194 - val\_loss: 0.0196 - val\_acc: 0.0184  
Epoch 25/150  
0s - loss: 0.0311 - acc: 0.0208 - val\_loss: 0.0196 - val\_acc: 0.0184  
Epoch 26/150  
0s - loss: 0.0296 - acc: 0.0222 - val\_loss: 0.0195 - val\_acc: 0.0184  
Epoch 27/150  
0s - loss: 0.0284 - acc: 0.0237 - val\_loss: 0.0194 - val\_acc: 0.0184  
Epoch 28/150  
0s - loss: 0.0291 - acc: 0.0186 - val\_loss: 0.0194 - val\_acc: 0.0184  
Epoch 29/150  
0s - loss: 0.0295 - acc: 0.0237 - val\_loss: 0.0193 - val\_acc: 0.0184  
Epoch 30/150  
0s - loss: 0.0288 - acc: 0.0201 - val\_loss: 0.0194 - val\_acc: 0.0184  
Epoch 31/150  
0s - loss: 0.0290 - acc: 0.0201 - val\_loss: 0.0193 - val\_acc: 0.0184  
Epoch 32/150  
0s - loss: 0.0283 - acc: 0.0229 - val\_loss: 0.0193 - val\_acc: 0.0184  
Epoch 33/150  
0s - loss: 0.0290 - acc: 0.0229 - val\_loss: 0.0192 - val\_acc: 0.0184  
Epoch 34/150  
0s - loss: 0.0297 - acc: 0.0222 - val\_loss: 0.0193 - val\_acc: 0.0184  
Epoch 35/150  
0s - loss: 0.0286 - acc: 0.0229 - val\_loss: 0.0192 - val\_acc: 0.0184  
Epoch 36/150  
0s - loss: 0.0278 - acc: 0.0237 - val\_loss: 0.0192 - val\_acc: 0.0184  
Epoch 37/150  
0s - loss: 0.0280 - acc: 0.0222 - val\_loss: 0.0193 - val\_acc: 0.0184  
Epoch 38/150  
0s - loss: 0.0310 - acc: 0.0194 - val\_loss: 0.0192 - val\_acc: 0.0184  
Epoch 39/150  
0s - loss: 0.0278 - acc: 0.0186 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 40/150  
0s - loss: 0.0286 - acc: 0.0201 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 41/150  
0s - loss: 0.0277 - acc: 0.0215 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 42/150  
0s - loss: 0.0275 - acc: 0.0222 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 43/150  
0s - loss: 0.0276 - acc: 0.0201 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 44/150  
0s - loss: 0.0273 - acc: 0.0222 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 45/150

0s - loss: 0.0285 - acc: 0.0229 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 46/150  
0s - loss: 0.0276 - acc: 0.0208 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 47/150  
0s - loss: 0.0285 - acc: 0.0194 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 48/150  
0s - loss: 0.0265 - acc: 0.0237 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 49/150  
0s - loss: 0.0282 - acc: 0.0201 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 50/150  
0s - loss: 0.0282 - acc: 0.0201 - val\_loss: 0.0191 - val\_acc: 0.0184  
Epoch 51/150  
0s - loss: 0.0280 - acc: 0.0215 - val\_loss: 0.0189 - val\_acc: 0.0184  
Epoch 52/150  
0s - loss: 0.0267 - acc: 0.0229 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 53/150  
0s - loss: 0.0286 - acc: 0.0215 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 54/150  
0s - loss: 0.0270 - acc: 0.0208 - val\_loss: 0.0189 - val\_acc: 0.0184  
Epoch 55/150  
0s - loss: 0.0271 - acc: 0.0208 - val\_loss: 0.0190 - val\_acc: 0.0184  
Epoch 56/150  
0s - loss: 0.0274 - acc: 0.0229 - val\_loss: 0.0189 - val\_acc: 0.0184  
Epoch 57/150  
0s - loss: 0.0267 - acc: 0.0222 - val\_loss: 0.0189 - val\_acc: 0.0184  
Epoch 58/150  
0s - loss: 0.0278 - acc: 0.0222 - val\_loss: 0.0188 - val\_acc: 0.0184  
Epoch 59/150  
0s - loss: 0.0256 - acc: 0.0201 - val\_loss: 0.0188 - val\_acc: 0.0184  
Epoch 60/150  
0s - loss: 0.0260 - acc: 0.0237 - val\_loss: 0.0188 - val\_acc: 0.0184  
Epoch 61/150  
0s - loss: 0.0276 - acc: 0.0201 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 62/150  
0s - loss: 0.0284 - acc: 0.0215 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 63/150  
0s - loss: 0.0268 - acc: 0.0229 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 64/150  
0s - loss: 0.0274 - acc: 0.0215 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 65/150  
0s - loss: 0.0264 - acc: 0.0215 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 66/150  
0s - loss: 0.0266 - acc: 0.0229 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 67/150  
0s - loss: 0.0261 - acc: 0.0229 - val\_loss: 0.0187 - val\_acc: 0.0184  
Epoch 68/150  
0s - loss: 0.0278 - acc: 0.0208 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 69/150

0s - loss: 0.0251 - acc: 0.0244 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 70/150  
0s - loss: 0.0275 - acc: 0.0222 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 71/150  
0s - loss: 0.0270 - acc: 0.0237 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 72/150  
0s - loss: 0.0274 - acc: 0.0244 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 73/150  
0s - loss: 0.0264 - acc: 0.0208 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 74/150  
0s - loss: 0.0261 - acc: 0.0244 - val\_loss: 0.0186 - val\_acc: 0.0184  
Epoch 75/150  
0s - loss: 0.0262 - acc: 0.0229 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 76/150  
0s - loss: 0.0260 - acc: 0.0222 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 77/150  
0s - loss: 0.0252 - acc: 0.0237 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 78/150  
0s - loss: 0.0259 - acc: 0.0201 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 79/150  
0s - loss: 0.0271 - acc: 0.0208 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 80/150  
0s - loss: 0.0260 - acc: 0.0222 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 81/150  
0s - loss: 0.0248 - acc: 0.0229 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 82/150  
0s - loss: 0.0278 - acc: 0.0222 - val\_loss: 0.0185 - val\_acc: 0.0184  
Epoch 83/150  
0s - loss: 0.0258 - acc: 0.0215 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 84/150  
0s - loss: 0.0260 - acc: 0.0229 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 85/150  
0s - loss: 0.0256 - acc: 0.0237 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 86/150  
0s - loss: 0.0244 - acc: 0.0237 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 87/150  
0s - loss: 0.0264 - acc: 0.0222 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 88/150  
0s - loss: 0.0256 - acc: 0.0244 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 89/150  
0s - loss: 0.0248 - acc: 0.0229 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 90/150  
0s - loss: 0.0274 - acc: 0.0229 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 91/150  
0s - loss: 0.0260 - acc: 0.0229 - val\_loss: 0.0184 - val\_acc: 0.0184  
Epoch 92/150  
0s - loss: 0.0252 - acc: 0.0229 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 93/150

0s - loss: 0.0247 - acc: 0.0194 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 94/150  
0s - loss: 0.0265 - acc: 0.0208 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 95/150  
0s - loss: 0.0250 - acc: 0.0222 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 96/150  
0s - loss: 0.0255 - acc: 0.0194 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 97/150  
0s - loss: 0.0258 - acc: 0.0237 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 98/150  
0s - loss: 0.0259 - acc: 0.0194 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 99/150  
0s - loss: 0.0254 - acc: 0.0222 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 100/150  
0s - loss: 0.0240 - acc: 0.0237 - val\_loss: 0.0183 - val\_acc: 0.0184  
Epoch 101/150  
0s - loss: 0.0251 - acc: 0.0208 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 102/150  
0s - loss: 0.0256 - acc: 0.0237 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 103/150  
0s - loss: 0.0264 - acc: 0.0215 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 104/150  
0s - loss: 0.0248 - acc: 0.0244 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 105/150  
0s - loss: 0.0258 - acc: 0.0237 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 106/150  
0s - loss: 0.0259 - acc: 0.0244 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 107/150  
0s - loss: 0.0254 - acc: 0.0251 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 108/150  
0s - loss: 0.0261 - acc: 0.0229 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 109/150  
0s - loss: 0.0239 - acc: 0.0222 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 110/150  
0s - loss: 0.0248 - acc: 0.0251 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 111/150  
0s - loss: 0.0243 - acc: 0.0237 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 112/150  
0s - loss: 0.0257 - acc: 0.0229 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 113/150  
0s - loss: 0.0246 - acc: 0.0215 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 114/150  
0s - loss: 0.0246 - acc: 0.0251 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 115/150  
0s - loss: 0.0250 - acc: 0.0237 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 116/150  
0s - loss: 0.0264 - acc: 0.0229 - val\_loss: 0.0182 - val\_acc: 0.0184  
Epoch 117/150

0s - loss: 0.0258 - acc: 0.0237 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 118/150  
0s - loss: 0.0261 - acc: 0.0201 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 119/150  
0s - loss: 0.0254 - acc: 0.0237 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 120/150  
0s - loss: 0.0253 - acc: 0.0244 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 121/150  
0s - loss: 0.0254 - acc: 0.0215 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 122/150  
0s - loss: 0.0253 - acc: 0.0244 - val\_loss: 0.0181 - val\_acc: 0.0184  
Epoch 123/150  
0s - loss: 0.0234 - acc: 0.0251 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 124/150  
0s - loss: 0.0239 - acc: 0.0237 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 125/150  
0s - loss: 0.0242 - acc: 0.0215 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 126/150  
0s - loss: 0.0243 - acc: 0.0222 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 127/150  
0s - loss: 0.0251 - acc: 0.0208 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 128/150  
0s - loss: 0.0264 - acc: 0.0222 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 129/150  
0s - loss: 0.0235 - acc: 0.0222 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 130/150  
0s - loss: 0.0255 - acc: 0.0215 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 131/150  
0s - loss: 0.0262 - acc: 0.0229 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 132/150  
0s - loss: 0.0249 - acc: 0.0237 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 133/150  
0s - loss: 0.0247 - acc: 0.0222 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 134/150  
0s - loss: 0.0266 - acc: 0.0208 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 135/150  
0s - loss: 0.0249 - acc: 0.0237 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 136/150  
0s - loss: 0.0265 - acc: 0.0215 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 137/150  
0s - loss: 0.0240 - acc: 0.0229 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 138/150  
0s - loss: 0.0241 - acc: 0.0251 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 139/150  
0s - loss: 0.0236 - acc: 0.0237 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 140/150  
0s - loss: 0.0261 - acc: 0.0222 - val\_loss: 0.0180 - val\_acc: 0.0184  
Epoch 141/150

```

0s - loss: 0.0255 - acc: 0.0229 - val_loss: 0.0180 - val_acc: 0.0184
Epoch 142/150
0s - loss: 0.0235 - acc: 0.0222 - val_loss: 0.0180 - val_acc: 0.0184
Epoch 143/150
0s - loss: 0.0248 - acc: 0.0229 - val_loss: 0.0179 - val_acc: 0.0184
Epoch 144/150
0s - loss: 0.0244 - acc: 0.0244 - val_loss: 0.0180 - val_acc: 0.0184
Epoch 145/150
0s - loss: 0.0249 - acc: 0.0229 - val_loss: 0.0179 - val_acc: 0.0184
Epoch 146/150
0s - loss: 0.0234 - acc: 0.0244 - val_loss: 0.0179 - val_acc: 0.0184
Epoch 147/150
0s - loss: 0.0231 - acc: 0.0251 - val_loss: 0.0179 - val_acc: 0.0184
Epoch 148/150
0s - loss: 0.0235 - acc: 0.0244 - val_loss: 0.0179 - val_acc: 0.0184
Epoch 149/150
0s - loss: 0.0232 - acc: 0.0229 - val_loss: 0.0179 - val_acc: 0.0184
Epoch 150/150
0s - loss: 0.0243 - acc: 0.0222 - val_loss: 0.0179 - val_acc: 0.0184
RMSE 0.22
Variance score: -0.01

```

**FNN doesnt perform as well possibly due to the limited number of training samples.**

```

In [33]: HighCorrTuples = []
        LowCorrTuples=[]
        Col = []
        ht = {}
        CorrCopy = copy.deepcopy(CorrList)
        for var in CorrCopy:
            if((var[0]>=0.8 or var[0]<=-0.8) and var[1]!=var[3]):
                if(var[1] not in ht or var[3] not in ht):
                    print var
                    ht[var[3]] =0
            else:
                LowCorrTuples.append(var)
        print len(ht)
        FinalFeatureSet = []
        print len(CorrCopy)
        for index,var in enumerate(CorrCopy):
            if(not(var[1] in ht or var[3] in ht)):
                #print var
                FinalFeatureSet.append(var)
        print len(CorrList)
        print len(CorrCopy)
        print len(FinalFeatureSet)
        '''

```



```

assert(len(HighCorrTuples)+len(LowCorrTuples)==len(Corr))
UncorrelatedFeatureIndices=[]

for i in HighCorrTuples:
    UncorrelatedFeatureIndices.remove(i[1])
print len(UncorrelatedFeatureIndices)

for i in LowCorrTuples:
    if(i[1]!=i[3]):
        if(i[1] not in UncorrelatedFeatureIndices):
            UncorrelatedFeatureIndices.append(i[1])
        if(i[3] not in UncorrelatedFeatureIndices):
            UncorrelatedFeatureIndices.append(i[3])
print len(UncorrelatedFeatureIndices)
print UncorrelatedFeatureIndices

'''

```

```

(0.985358028188035, 47, 'v_cont_49', 46, 'v_cont_48', 0.0)
(0.9311545163340417, 47, 'v_cont_49', 48, 'v_cont_50', 0.0)
(0.9076518065031174, 47, 'v_cont_49', 49, 'v_cont_51', 0.0)
(-0.8701983204454441, 47, 'v_cont_49', 53, 'v_cont_55', 0.0)
(0.985358028188035, 46, 'v_cont_48', 47, 'v_cont_49', 0.0)
(-0.8034519072989161, 6, 'v_cont_8', 53, 'v_cont_55', 0.0)
(0.8132433874202087, 18, 'v_cont_20', 22, 'v_cont_24', 0.0)
(0.981898517014815, 70, 'v_cont_72', 76, 'v_cont_78', 0.0)
(0.981898517014815, 76, 'v_cont_78', 70, 'v_cont_72', 0.0)
(0.8132433874202087, 22, 'v_cont_24', 18, 'v_cont_20', 0.0)
(0.9776158180285813, 22, 'v_cont_24', 15, 'v_cont_17', 0.0)
(0.9438503379649609, 32, 'v_cont_34', 33, 'v_cont_35', 0.0)
(0.8814456236410259, 74, 'v_cont_76', 30, 'v_cont_32', 0.0)
(0.8226685567453994, 74, 'v_cont_76', 52, 'v_cont_54', 0.0)
(0.8814456236410259, 30, 'v_cont_32', 74, 'v_cont_76', 0.0)
(0.9438503379649609, 33, 'v_cont_35', 32, 'v_cont_34', 0.0)
(0.8137998648445638, 80, 'v_cont_82', 31, 'v_cont_33', 0.0)
(0.8361952961320445, 35, 'v_cont_37', 20, 'v_cont_22', 0.0)
(0.8137998648445638, 31, 'v_cont_33', 80, 'v_cont_82', 0.0)
(0.9757255165940261, 41, 'v_cont_43', 44, 'v_cont_46', 0.0)
(0.9223946342691782, 41, 'v_cont_43', 43, 'v_cont_45', 0.0)
(0.9757255165940261, 44, 'v_cont_46', 41, 'v_cont_43', 0.0)
(0.8361952961320445, 20, 'v_cont_22', 35, 'v_cont_37', 0.0)
(0.8107003756835877, 5, 'v_cont_7', 53, 'v_cont_55', 0.0)
(-0.8034519072989161, 53, 'v_cont_55', 6, 'v_cont_8', 0.0)
(0.8107003756835877, 53, 'v_cont_55', 5, 'v_cont_7', 0.0)
24
729
729
729

```

9


```
Out[33]: '\nassert(len(HighCorrTuples)+len(LowCorrTuples)==len(Corr))\nUncorrelatedFeatureIndi
```

```
In [34]: Indfor ={}
         for ind,x in enumerate(FinalFeatureSet):
             if(x[1] not in Indfor):
                 Indfor[x[1]] = x[1]
         print len(Indfor)
         IndforCor =[]
         for key,val in Indfor.items():
             IndforCor.append(val)
         Cols=[]
         FinalcorMatrix = np.zeros((len(Indfor),len(Indfor)))
         for i in range(FinalcorMatrix.shape[0]):
             Cols.append(X.columns[IndforCor[i]])
             for j in range(FinalcorMatrix.shape[1]):
                 t1,t2 = pearsonr(Matrix[:,IndforCor[i]],Matrix[:,IndforCor[j]])
                 FinalcorMatrix[i][j] = t1

         fig, ax = plt.subplots(figsize=(20,20))
         sb.heatmap(FinalcorMatrix,
                     xticklabels=Cols,
                     yticklabels=Cols,linewidths=1,annot=True, ax=ax,square =True)

         plt.show()
```

3



```
In [35]: print IndforCor
[72, 77, 71]
```

### 1.5.3 Linear regression on the remaining features

```
In [36]: FinalTrainSetForTrain = np.zeros((3,Matrix.shape[0]))
         for ind,k in enumerate(IndforCor):
             FinalTrainSetForTrain[ind] = (Matrix[:,k])
         FinalTrainSetForTrain = FinalTrainSetForTrain.T
         X_train,X_test,Y_train,Y_test = train_test_split(FinalTrainSetForTrain,Y,test_size = 0.2)
```

```
#Normal Linear Regression
```

```
clf = linear_model.LinearRegression(normalize = True)
clf.fit(X_train,Y_train)
Y_pred = clf.predict(X_test)
Weights = clf.coef_
print("RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred)))
print('Variance score: %.2f' % r2_score(Y_test, Y_pred))
```

```
# Linear Regression With Ridge regularisation
```

```
clf = linear_model.Ridge(alpha = 0.1,normalize = True)
clf.fit(X_train,Y_train)
Y_pred = clf.predict(X_test)
print("RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred)))
print('Variance score: %.2f' % r2_score(Y_test, Y_pred))
```

```
# Linear Regression With Lasso regularisation
```

```
clf = linear_model.Lasso(alpha = 0.1,normalize = True)
clf.fit(X_train,Y_train)
Y_pred = clf.predict(X_test)
print("RMSE %.2f" % math.sqrt(mean_squared_error(Y_test,Y_pred)))
print('Variance score: %.2f' % r2_score(Y_test, Y_pred))
```

```
RMSE 0.17
Variance score: 0.42
RMSE 0.17
Variance score: 0.42
RMSE 0.22
Variance score: -0.01
```